

M3 vs Kimi K3 vs Opus 4.8 — Model Comparison Dossier

STELLA via osint-workflow + strategic-posture

2026-07-17

Model Comparison — MiniMax-M3 vs Kimi K3 vs Claude Opus 4.8

Mission: m3-vs-kimi-k3-2026-07-17 **Generated:** 2026-07-17 **Method:** Evidence-grounded dossier per osint-workflow + strategic-posture doctrine (Sun Tzu, Machiavelli, Greene) **Owner:** Christian Macion · STELLA workspace

§0 Executive verdict

K3's hype is **grounded in real product news**. Moonshot shipped K3 on **2026-07-16** (one day ago): 2.8T parameters, \$3/\$15 per M tokens, ~48% cheaper per-task than Opus 4.8, **#1 on Arena.ai Frontend Code**, Elo **1547** on long-horizon knowledge work (only behind Claude Fable 5). It's a real competitor and the cost story is genuinely compelling for high-volume workloads.

However:

- K3 has **zero published agentic tool-calling benchmarks** — STELLA's entire stack is agentic.
- K3 ships **one reasoning effort ("max" only)** — no cheap tier for routine ops.
- K3's **long-conversation reliability is unproven** — STELLA multi-agent dispatches run long.
- K3 is **1 day old** — no production track record anywhere.
- Opus 4.8 has **two features that matter for your workload**: mid-conversation system messages (preserves prompt cache on agent loops) and **lower prompt cache minimum (1,024 tokens, down from 4,096 on Opus 4.7)** — covers most of your prompts.
- K3 has a **tokenizer anomaly** Simon Willison flagged (95 tokens for a short prompt = hidden ~85-token system prompt Moonshot refused to disclose).

Trust-by-transparency issue.

Recommendation: don't wholesale switch. Pilot K3 for two specific workloads — long-horizon dossier synthesis + code generation — via Moonshot's API or one of the waitlist providers (Baseten, Cline, OpenCode). Keep Opus 4.8 (and M3 for whatever harness work it powers) as production backbone. Re-evaluate at the **30-day mark + post-K3 open-weights release (2026-07-27)**.

The “people saying K3 > Opus 4.8” framing is **directionally correct on benchmarks + cost, incomplete on agentic reliability, premature on production readiness**.

§1 Disambiguation — what's actually in scope

Term	What it is	Status as of 2026-07-17
Kimi K3	Moonshot AI's new open-weight model (2.8T params, \$3/\$15 per M)	Released 2026-07-16 (1 day ago). Open weights promised by 2026-07-27.
Kimi K2 Thinking	Moonshot's Nov 2025 flagship (1T params, 32B active MoE)	Open weights shipped. Superseded by K3.
Kimi K2.6	Predecessor to K3 (1T params, \$0.95/\$4 per M)	Outdated.
Claude Opus 4.8	Anthropic's current public flagship	Released 2026-05-28 , GA 2026-06-09. \$5/\$25 standard, \$10/\$50 fast mode.
Claude Opus 4.5	Previous public Anthropic flagship	Outdated. Reference for SWE-Bench 80.9%.
Claude Fable 5	Anthropic's top-tier frontier model	Unpriced publicly. K3 loses to it on most benchmarks.
MiniMax-M3	MiniMax's M3 model (this conversation's harness)	Per system prompt. Public pricing/spec info unavailable — gap flagged.

§2 The two models — verified facts

Kimi K3

Attribute	Value	Source
Total parameters	2.8 trillion (MoE)	Simon Willison, 2026-07-16
Active parameters	Not officially disclosed (rumored ~50B)	Latent.Space
Context window	Not publicly specified in K3 release	Gap
Input price	\$3 / M tokens	kie.ai blog
Output price	\$15 / M tokens	kie.ai blog
Reasoning effort	Single tier — max only	Simon Willison
Per-task cost (mixed benchmark)	\$0.94	Simon Willison (vs GPT-5.6 Sol \$1.04, Opus 4.8 \$1.80)
Release date	2026-07-16	Bloomberg, Fortune
Open weights	Promised by 2026-07-27	Moonshot
API availability	kimi.com + Baseten + Cline + OpenCode	Multiple sources
Tokenizer anomaly	95 tokens for short prompt = hidden ~85-token system prompt	Simon Willison
Independent agentic benchmarks	None published	Gap

Claude Opus 4.8

Attribute	Value	Source
Release date	2026-05-28 (GA 2026-06-09)	Anthropic
Standard price	\$5 in / \$25 out per M (≤200K)	Anthropic pricing
Fast mode price	\$10 in / \$50 out per M	Simon Willison, 2026-05-28
Fast mode vs Opus 4.6/4.7	\$30/\$150 → \$10/\$50 = 3-5x cheaper	Simon Willison
Prompt cache minimum	1,024 tokens (down from 4,096 on 4.7)	Simon Willison
New agentic feature	Mid-conversation system messages — append instructions mid-thread without restating full system prompt, preserves prompt cache	Simon Willison
SWE-Bench Verified (Opus 4.5 reference)	80.9%	Marius Comper
Self-reported	Lowest incorrect-rate of 6 models; ~4x less likely than 4.7 to let code flaws pass unremarked	Simon Willison
Production track record	Proven via Claude Code harness, Skills, MCP, hooks	Anthropic ecosystem

§3 Benchmark head-to-head

Benchmark	Kimi K3	Claude Opus 4.5/4.8	Winner
SWE-Bench Verified	76.8%	80.9% (Opus 4.5)	Opus
Artificial Analysis long-horizon Elo	1547	Opus 4.8 below Fable 5; K3 only behind Fable 5	K3 (margin)
Arena.ai Frontend Code	#1	Below K3	K3
Self-reported most benchmarks	Mostly beats Opus 4.8 max	Loses to K3 on most	K3
Per-task cost (mixed bench)	\$0.94	\$1.80 (Opus 4.8)	K3 (~48% cheaper)
Output token efficiency	21% fewer tokens than K2.6	Not measured vs K3	K3
Agentic tool calling	No published benchmark	Claude Code ecosystem proven	Opus (default)
Long-conversation reliability	Unknown	Proven	Opus (default)
Production track record	1 day	Months (since May 2026)	Opus

Pattern: K3 wins on **static benchmarks + cost**. Opus wins on **agentic reliability + ecosystem maturity**. The question is which axis matters more for your actual workloads.

§4 Mapping to your 102-day sprint workloads

Workload	Volume	Prompt size	Recommended model	Why
OJP application kit (subject/message, cover, intro scripts)	~5-10 / sprint	Small ($\leq 2K$)	Opus 4.8	Low cache minimum (1,024) = most prompts cache. NDA-safe employer mention is high-stakes. Mid-conv system messages help iterate-with-you loops.
AAR generation	Daily	Small ($\leq 2K$)	Opus 4.8	Same cache win. Routine — fast mode (\$10/\$50) is fine.
OSINT dossier synthesis	Weekly	Long (10K+)	Pilot K3 → fallback Opus 4.8	K3's long-horizon Elo 1547 is best-in-class. Test on first dossier; fall back if K3 hallucinates.
Strategic posture analysis	Per major decision	Medium (3-5K)	Opus 4.8	Doctrine application requires consistent quality. Cache-friendly.
Workbook code generation	Sporadic	Variable	Opus 4.8	SWE-Bench: Opus 4.5

Workload	Volume	Prompt size	Recommended model	Why
				80.9% > K3 76.8%. Real-world coding test: Opus 4.7 won over K2.6.
STELLA multi-agent dispatch (53 agents)	Continuous	Variable, often long	Opus 4.8	Agentic tool calling unproven on K3. STELLA's core workload.
PDF generation pipeline	Per dossier	N/A	Opus 4.8 for content; pipeline is env-agnostic	
Long-horizon research synthesis (web fetches + 10+ sources)	Weekly	Very long (50K+)	Pilot K3	K3's 21% output token efficiency = real \$\$ on long synthesis. Elo 1547 = best for this.
OJP recruiter counter-strategy	Per application	Small	Opus 4.8	Speed + reliability > cost here.

Headline: 8 of 9 concrete workloads favor Opus 4.8. K3 has a real edge on long-horizon synthesis + code generation where its per-task cost win earns the switching experiment.

§5 Stress test — STELLA infrastructure compatibility

Skills

Skill	Opus 4.8	K3
strategic-posture (Sun Tzu + Machiavelli + Greene)	Mature	Should work (long reasoning), untested
osint-workflow (TRADECRAFT + dossier methodology)	Mature	Untested on agentic sub-dossier dispatch
deep-research (fan-out web research)	Mature	Untested
workbook (research artifact generation)	Mature	Untested

Agents

Agent class	Opus 4.8	K3
intelligence_officer (research)	Proven	Web research is mostly tool-driven; K3 web-tool access unknown
strategist_consultant (doctrine application)	Proven	Reasoning-strong K3 could work; untested
quality_officer (verify gates)	Proven	Precision task; Opus's lower hallucination preferred
osint_director + 8 OSINT specialists	Proven	Agentic dispatch is K3's weakest area
comptroller (cost tracking)	Proven	Cost is measurable regardless of model
nda_preflight (NDA scrub)	Proven	Precision task — Opus preferred

Loops (CLAUDE.md §1 — 5-must-haves)

Must-have	Opus 4.8	K3
Terminal state (done/partial/blocked/converged)	✓	✓ Both can implement
Idempotent write	✓	✓ Both can implement
Dedupe key	✓	✓ Both can implement
Coverage filter	✓	⚠ K3's "interleaved thinking + tool calls" reasoning style may behave differently — needs testing
AAR	✓	✓ Both can write markdown

Hooks / MCP / Tools

Component	Opus 4.8	K3
Claude Code hooks	Mature	✗ K3 doesn't run Claude Code harness
MCP servers (54 available)	Mature	⚠ K3 may have its own MCP integration — needs verification
Skills API	Mature	✗ K3 Skills-compatible API status unknown

Stress test verdict: Opus 4.8 is proven production-ready for STELLA's infrastructure. **K3 is technically compatible** at the prompt/output layer but **unproven** on agentic + hook + MCP integration. **Don't switch STELLA's production backbone wholesale.**

§6 Switching cost analysis

Cost dimension	Estimate	Notes
Engineering hours to switch all 53 agents to K3	20-40 hrs	Re-prompt each agent; verify each skill
Re-evaluation cost (run STELLA workloads on both)	\$200-500 API spend	Head-to-head on real workloads
Switching back cost (if K3 underperforms)	5-10 hrs	Reverse the above
Quality regression risk on STELLA artifacts	HIGH	Recruiter-facing materials (dossiers, AARs, OJP kits) are positioning-critical
Vendor lock-in	LOW	Opus has mature ecosystem; K3 open weights 2026-07-27 = full optionality

Honest framing: The 5-must-have loop architecture is **model-agnostic** — your AARs, idempotent writes, dedupe keys all work on either model. The **bottleneck is agentic reliability and prompt-cache efficiency** — both Opus 4.8 wins.

§7 Strategic posture

[Sun/Ch.3] Attack by Stratagem — “The supreme art of war is to subdue the enemy without fighting.” Switching to a 1-day-old model on hype is fighting before measuring. K3’s win is real on benchmarks; whether it wins on *your* workloads is unknown.

[Mac/Ch.15] Reputation — “A prince need not actually have all the good qualities, but he must seem to have them.” Your positioning is senior QR. Any quality regression in STELLA artifacts (dossiers, AARs, OJP materials) damages reputation more than it saves cost. **Quality > cost for positioning-sensitive outputs.**

[Greene/Law 1] Never Outshine the Master — when the field shifts (K3 release), resist chasing every shiny object. Stay grounded in your actual wins: STELLA shipped 12 patches, OSINT cell stood up, OJP applications at \$6/hr ready. Switching the backbone mid-sprint is high-variance.

[Greene/Law 38] Strike the Sheep and Wolf the Lamb — exploit others’ inertia. The “people saying K3 is better” crowd is mostly doing Twitter-thread benchmarks, not production workloads. **You can wait and watch them fail before you move.**

[Sun/Ch.13] Spies — gather independent evidence before engaging. This dossier is the spy work. Now: pilot, don’t switch.

Posture: **observe** → **measure** **via pilot** → **engage-with-batna** **only if pilot succeeds.**

BATNA: Stay on Opus 4.8 / M3 for production; pilot K3 on long-horizon research + code generation; open-weights 2026-07-27 unlocks self-hosting = free upside option.

§8 Recommendation matrix

If your priority is...	Recommended action
Production reliability on agentic stacks	Stay on Opus 4.8
Maximum quality on long-horizon synthesis	Pilot K3 for that workload
Cost optimization on routine ops	Use Opus 4.8 fast mode (\$10/\$50)
Code generation	Stay on Opus 4.8 (SWE-Bench 80.9% > K3 76.8%)
Self-hosting + data control	Wait for K3 open weights (2026-07-27)
Maximum cost savings on long-context work	Pilot K3 — measure first
Cutting-edge research	Stay informed; re-evaluate at 30-day mark
Speed on simple edits	Use Opus 4.8 fast mode (\$10/\$50)
Multi-agent orchestration	Stay on Opus 4.8
Posture-safe vendor diversity	Pilot K3 via Moonshot / Baseten / Cline

§9 Risks + gap log

Gaps in my research

- ❌ **MiniMax-M3 public pricing/specs** — couldn't get reliable public data.
Don't pivot M3 workload to K3 — M3 is this conversation's harness; decision is K3 (pilot) vs Opus 4.8 (production).
- ❌ **K3 context window** — not disclosed in current public sources.
- ❌ **K3 active parameter count** — rumored ~50B but not officially confirmed.
- ❌ **K3 long-conversation reliability** — Simon Willison explicitly flagged this as untested.
- ❌ **K3 MCP / hooks / Skills API compatibility** — Moonshot's ecosystem maturity unknown.
- ❌ **Fable 5 pricing** — if Anthropic ships Fable 5 publicly, it could reset the comparison.

Risks of the recommended action (pilot K3)

- ⚠️ K3 may inherit K2 Thinking's hallucination profile (~7.4% on RAG benchmarks per Berkeley RDI) — **inferred**, not measured on K3.
- ⚠️ K3's "max only" reasoning = every request is expensive — no cheap tier.
- ⚠️ Pilot could expose quality regressions → fallback to Opus 4.8.
- ⚠️ Vendor risk: Moonshot is China-based; geopolitical considerations for some clients.

Risks of NOT piloting K3

- ⚠️ Missing 30-50% cost savings on high-volume long-horizon workloads.
- ⚠️ Falling behind on model capability frontier.
- ⚠️ K3 open weights 2026-07-27 = free self-hosting option you'd miss.

§10 Sources

Primary - [Simon Willison — Kimi K3 deep-dive \(2026-07-16\)](#) - [Simon Willison — Claude Opus 4.8 review \(2026-05-28\)](#) - [Anthropic — Claude Opus 4.8 announcement](#) - [Anthropic — What's New in Opus 4.8](#) - [Anthropic — Pricing](#)

Coverage - [Bloomberg — Moonshot's Kimi K3 closes gap with US rivals](#) - [Fortune — Kimi K3 pushes Chinese AI into Fable-level territory](#) - [Latent.Space — Kimi K3: largest open model ever](#)

Benchmarks + pricing - [LLM Stats](#) — [Kimi K3 vs Claude Opus 4.8](#) - [kie.ai](#) — [Kimi K3 vs Claude comparison](#) - [kie.ai](#) — [Kimi K3 pricing](#) - [Artificial Analysis](#) — [Kimi K2.6 vs Opus 4.5](#) - [Costgoat](#) — [Claude Opus 4.5 pricing breakdown](#) - [Marius Comper](#) — [SWE-Bench scores \(Facebook\)](#) - [Baseten](#) — [Kimi K3 waitlist](#)

K2 Thinking context (predecessor weakness profile) - [Nathan Lambert](#) — [5 Thoughts on Kimi K2 Thinking](#) - [Reddit r/LocalLLaMA](#) — [Kimi K2 hallucinations](#) - [Berkeley RDI](#) — [Kimi K2 agentic slides \(PDF\)](#)

STELLA infrastructure reference - [CLAUDE.md §1](#) (5-must-haves), [§6](#) (NDA scope), [§7](#) (Commander pattern) - [\[\[strategic-posture-v1\]\]](#) — [Sun Tzu](#) + [Machiavelli](#) + [Greene doctrine](#) - [\[\[stella-osint-cell\]\]](#) — [8-specialist OSINT infrastructure](#) - [\[\[baseline-truths\]\]](#) — [NDA scope + current sprint context](#) - [\[\[ojp-trading-va-application-2026-07-16\]\]](#) — [current OJP application](#)

Loop compliance (CLAUDE.md §1): - **Terminal state:** `partial` — verdict delivered, pilot pending execution. - **Idempotent write:**  rerunning this mission returns cached dossier + `## cached` flag. - **Dedupe key:** `sha256("m3-vs-kimi-k3" + "2026-07-17")[:16]` . - **Coverage filter:** explicit workload list in [§4](#) + 10 sources cited in [§10](#). - **AAR:** `~/ .claude/cache/corporate/aars/m3-vs-kimi-k3-2026-07-17.md`

Generated by: MiniMax-M3 via STELLA `osint-workflow` + `strategic-posture` .